

① Introduction

- How to sync content across systems?

Content

1. Problem formulation
2. Gossip protocols
3. Network coding
4. Fountain codes
5. De-sync

① Problem:

- Before: Transferring files reliably
but if we want synchronize content across systems
- Real problem example:
"XEROX": maintain a distributed database, all replicas should be coherent (latest and consistent data)

② Initial Approaches

1. Initial:
 - a) Local Updates: updates happened locally
 - b) Update Propagation
 - c) Method: spread via email
 - d) Problems: #3

(run in a background to recover automatically)

2. Anti-Entropy: simple epidemics. [Randomly contact sites]

def: entropy - tendency of system to become increasingly disordered and equalized

- keeping system Ordered (low H) requires energy

- we have 2 types of nodes:

S - need 1 or more update and I - has at least one recent update

- if start with one I after $\log(n) + \ln(n) + O(1)$ (with push) we will affect the entire population using $O(n \log n)$ messages

- depending on the design we have 3 operations, in each we compare entire DBs:

(choose to contact)
1. Push 2. pull 3. push-pull . S - site remain S after i+1 cycle
 $p_{i+1} = p_i \left(1 - \frac{1}{n}\right)^{n(1-p_i)}$ $p_{i+1} = (p_i)^2$ -//-

- Problem: copy the entire DB just for small changes

3. Anti-entropy: Complex Epidemics Rumor spreading [Every small update = rumor]

- S - doesn't know about update, I - knows and shares, R - knows, not shares

- $S+i+r=1$
$$\frac{ds}{dt} = -si; \quad \frac{di}{dt} = si - \frac{1}{k}(1-s) \cdot i; \quad \frac{dr}{dt} = \frac{1}{k}(1-s) \cdot i$$

$$S = e^{-(k+1)(1-s)t}$$

- When to lose interest?

- after k unnecessary contacts

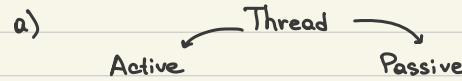
- after 1

- with $p = \frac{1}{k}$

- Some prob. of errors

Idea: check for updates using Anti-entropy, but spread small update sets with rumors

③ Gossip protocols



Active thread:
once in each T time units,
at a random time, do:
`p = selectPeer();`
send mystate to p
receive hisstate from p
mystate = update(hisstate)

Passive thread:
do forever:
receive herstate from p
send mystate to p
mystate = update(herstate)

Requirements:

1. Random peer selection

3. round-based communication

5. Same protocol for all

2. Nodes use only local info

4. limited transmission and processing per round

- b) Rumor mongering - special case
- c) \downarrow selectPeer() \downarrow
 full connectivity Only local actions
- d. Bootstrap:
1. „Sampling service“
 2. Wanna tuning?:
 T-MAN (protocol to manage and optimize the network topology)
- e. Gossip for Distributed Computations:
- Structured Aggregation Unstructured aggregation
- Scatter/Gather Algo
 - State = current approx. of result
- Random Select
 $\frac{s_1+s_2}{\sqrt{s_1 \cdot s_2}}$
 $\rightarrow \max(s_1, s_2)$

④ Network coding

1. Separation theorem
2. Allow intermediate nodes to perform coding op.
3. „Butterfly configuration“
4. Problems: topology-awareness + synchronized stream of inputs

⑤ Fountain Codes

1. Broadcasting problem
2. Approach "Data Carousel" + Raptor codes

⑥ De-sync Algo

② Identifiers and Network Association

- How devices and users identify themselves and other within a network?

① Problem

- Chicken-and-egg problem (paradox in establishing a shared view of the network)

② Introduction

- Divide and Conquer Approach:

Conventions on Identifying Scheme

↳ Examples:

- IP
 - IPv4 - 32 bit
 - IPv6 - 128 bit
- MAC (Assigned to NIC,
- Port service number
- DNS (Domain Name System)
- Structure
 - Hierarchical
 - flat
- more effective storage
- not scalable
- Ex.: DNS
 - Ex. MAC, UUIDs
- Mappings Identifier spaces:
 - DNS
 - LDAP
 - White pages

Conventions on Identifier Acquisition Mechanism

- Request- Allocate (Client request → network gives)
- Select - Advertise (Client selects → network)
 - ↳ Self-Certifying Names (hash, crypto keys)

Background knowledge:

③ Network association

- Because of node mobility, we need automatic identifier acquisition

Mechanisms:

① Dynamic Host Configuration Protocol (DHCP)

- def

- adv. + dis.

② IPv6 Stateless Address Auto Configuration (SLAAC)

- 1. Generate Link-Local address
2. Host sends Router Solicitation (RS) to "all routers"
3. Routers reply Router Advertisement Messages
4. Host generates an address in that prefix
5. Duplicate Address Detection (DAD)
6. Active address.

- Pros + cons (only for address auto-configuration)

• IPv6 Prefix Delegation via DHCPv6

- a network dynamically obtain a range of IPv6 addresses

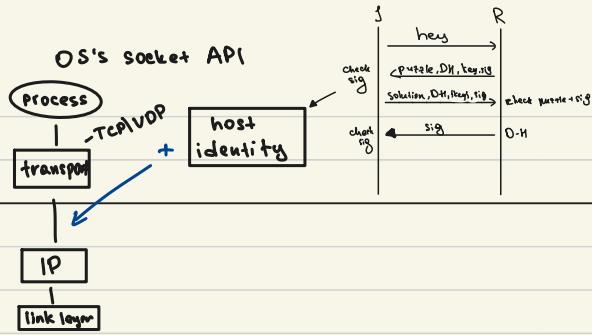
- Using 2 DHCPv6

③ IEEE 802.11 Radio Network Auto Association

④ Multiple identifier schemes in network association

- we can move and switch IPs, but we still we (→ privacy)

① HIP



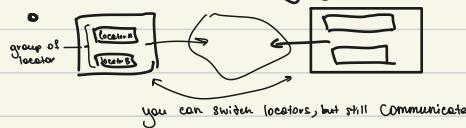
• LISP / Onions!

+ RVS (Rendezvous server)

② LISP (Overlay IP topology)

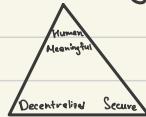
- Edge network \rightarrow IP = EID \rightarrow ITR (Ingress Tunnel Router)
- Core network \rightarrow IP = RLOC \rightarrow ETR (Egress Tunnel Router)

③ SHIM6 (Site Multi-Homing by IPv6 Intermediation)



⑤ Decentralized Identifiers (DID)

- Use public key
- Let's encrypt
- GNU Name System (GNS)
- Zooko's Triangle



③ - ④ Cryptography concepts

- How to communicate in private and know with whom?

① Negligible function

$\negl: \mathbb{N} \rightarrow \mathbb{R}$, for all $x' > x$

$$|\negl(x)| < \frac{1}{\text{poly}(x)}$$

• Asymptotic def (PPT-probabilistic polynomial-time):

A scheme is *secure* if for every PPT adversary \mathcal{A} carrying out an attack of some formally specified type, and for every positive polynomial p , there exists an integer N such that when $n > N$ the probability that \mathcal{A} succeeds in the attack is less than $\frac{1}{p(n)}$.

Note that nothing is guaranteed for values $n \leq N$.

② Adversary Models:

1. Semi-honest (follows, but curious)
2. Malicious (arbitrarily deviate)
3. Covert (act maliciously, min. risk of detection)
4. Rational (act based on self-interest)

3. Kerckhoff's Principal

Ex. $ax+b$ with a, b - secret

- Mechanisms must be secure, even if.

• The algo/protocol is public

• params are public

• only random bits used during computations ("keys") can be assumed secret

=> Privacy without obscurity

④ Symmetric Key Encryption

- Idea

• DEFINITION 3.7 A private-key encryption scheme is a tuple of probabilistic polynomial-time algorithms $(\text{Gen}, \text{Enc}, \text{Dec})$ such that:

1. The key-generation algorithm Gen takes as input 1^n (i.e., the security parameter written in unary) and outputs a key k ; we write $k \leftarrow \text{Gen}(1^n)$ (emphasizing that Gen is a randomized algorithm). We assume without loss of generality that any key k output by $\text{Gen}(1^n)$ satisfies $|k| \geq n$.

2. The encryption algorithm Enc takes as input a key k and a plaintext message $m \in \{0,1\}^*$, and outputs a ciphertext c . Since Enc may be randomized, we write this as $c \leftarrow \text{Enc}_k(m)$.

3. The decryption algorithm Dec takes as input a key k and a ciphertext c , and outputs a message m or an error. We assume that Dec is deterministic, and so write $m := \text{Dec}_k(c)$ (assuming here that Dec does not return an error). We denote a generic error by the symbol \perp .

It is required that for every n , every key k output by $\text{Gen}(1^n)$, and every $m \in \{0,1\}^*$, it holds that $\text{Dec}_k(\text{Enc}_k(m)) = m$.

- AES : todo

⑤ Hash function

- def: • Pair(Gen, H)

• Collision: $\text{find}: x, x' \text{ find } x': \text{H}(x) = \text{H}(x')$

• Target-collision: $s, x \text{ find } x': \text{H}(x) = \text{H}(x')$

• Preimage resistant: $s, y \text{ find } x: \text{H}(x) = y \leftarrow \text{one-way}$

• SHA256

⑥ Public key Encryption

- Idea

•

- Example: RSA, ElGamal todo

DEFINITION 11.1 A public-key encryption scheme is a triple of probabilistic polynomial-time algorithms $(\text{Gen}, \text{Enc}, \text{Dec})$ such that:

1. The key-generation algorithm Gen takes as input the security parameter 1^n and outputs a pair of keys (pk, sk) . We refer to the first of these as the public key and the second as the private key. We assume for convenience that pk and sk each has length at least n , and that n can be determined from pk, sk .

2. The encryption algorithm Enc takes as input a public key pk and a message m from some message space (that may depend on pk). It outputs a ciphertext c , and we write this as $c \leftarrow \text{Enc}_{pk}(m)$. (Looking ahead, Enc will need to be probabilistic to achieve meaningful security.)

3. The deterministic decryption algorithm Dec takes as input a private key sk and a ciphertext c , and outputs a message m or a special symbol \perp denoting failure. We write this as $m := \text{Dec}_{sk}(c)$.

It is required that, except possibly with negligible probability over (pk, sk) output by $\text{Gen}(1^n)$, we have $\text{Dec}_{sk}(\text{Enc}_{pk}(m)) = m$ for any (legal) message m .

7. Digital Signature

- Ex. DSA

8. Discrete Log Problem

- $g^x \bmod p = y$, G-group of prime order p with gen. g

9. D-H key exchange

- multiplicative group of int. mod N

number elements - order

- elements whose powers cycle through group - primitive roots mod N

- cyclic group - at least 1 generator

- shared: N-group, g

privat key: pick 1 number from N

$$S = g^{\text{privat}} \bmod N$$

$$S^{\text{privat}} \bmod N$$

$$B^a \bmod N = (g^b \bmod N)^a \bmod N = g^{ab} \bmod N$$

hard to undo modular exponentiation (discrete log problem)

10. Commitment scheme

- g, h - generators

$$g^m h^n = C$$

- (m, n) - commitment

11. Zero-knowledge Proof

- Schnorr Protocol ($h = g^a$)

Prover: choose a random $r \in G$

$$t = g^r$$

Verifier: send $c \in G$

Prover: $s = r + ca$

$$\text{Verifier: } g^s = g^{r+ca} = g^r \cdot g^{ca} = t \cdot h^c$$

$$\begin{array}{c} h = g^a \\ g^b \in G \\ \text{Verifier: } c \in G \\ \text{Prover: } \text{④ } b + h \cdot c \\ g^d = g^b \cdot g^c \end{array}$$

12. Multiparty computations

- Preknowledge.

- def Group((S, \circ)) - closure, associativity, identity, inverse)

- group homomorphism

$$(G, \circ) \text{ and } (H, \circ), f: G \rightarrow H, f(g \circ g') = f(g) \circ f(g')$$

- ring homom

$$h(x+y) = h(x) \otimes h(y) \text{ and } h(x \cdot y) = h(x) \otimes h(y); h: G \rightarrow H$$

① Partial Homo Encr:

- El Gamal (\circ)

$$\left\{ \begin{array}{l} \circ G, q, g \\ \circ x \in \{1, \dots, q-1\} - \text{randomly, private key} \end{array} \right.$$

- El Gamal (\circ)

$$\left\{ \begin{array}{l} -/- \\ -/- \end{array} \right.$$

G, g, q

$x \in \{1, \dots, q-1\} - \text{sk}, y \in \{1, \dots, q-1\}$

$$h = g^x \Rightarrow (G, q, g, h) - \text{pk}$$

$$C_1 = g^y \quad C_2 = m \cdot h^y \quad g^m \cdot h^y$$

$$\frac{C_2}{C_1^x} = \frac{m \cdot h^y}{(g^y)^x} = \frac{m \cdot g^{xy}}{g^{yx}}$$

$$\left\{ \begin{array}{l} \circ h = g^x \Rightarrow (G, q, g, h) - \text{public key} \end{array} \right.$$

- encryption

choose y from $\{1, \dots, q-1\}$

$$\circ C_1 = g^y \quad C_2 = m \cdot h^y = m \cdot (g^x)^y$$

Ciphertext $E(M) = (C_1, C_2)$

$$\left\{ \begin{array}{l} C_1 = g^y \\ C_2 = m \cdot h^y \\ E(M) = (C_1, C_2) \end{array} \right.$$

$$\left\{ \begin{array}{l} \frac{C_2}{C_1^x} = \frac{m \cdot h^y}{(g^y)^x} = \frac{m \cdot g^{xy}}{g^{yx}} \\ \text{then brut force for } m \end{array} \right.$$

$$\left\{ \begin{array}{l} \text{decryption} \\ m = \frac{c_2}{c_1^x} = \frac{m \cdot h^y}{g^{xy}} = \frac{m \cdot g^{xy}}{g^{xy}} = m \end{array} \right.$$

$$\begin{aligned} E(m_1) \cdot E(m_2) &= (g^{y_1}, m_1, h^{y_1}) \cdot (g^{y_2}, m_2, h^{y_2}) = (g^{y_1+y_2}, (m_1 \cdot m_2) h^{y_1+y_2}) = E(m_1 \cdot m_2) \\ E(m_1) + E(m_2) &= (g^{y_1}, g^{m_1}, h^{y_1}) + (g^{y_2}, g^{m_2}, h^{y_2}) = (g^{y_1+y_2}, g^{m_1+m_2}, h^{y_1+y_2}) \end{aligned}$$

② Somewhat HE Scheme

- s and e vectors of n random numbers in mod q
- A' - matrix $n \times N$ random numbers
- $b = A's + 2e$
- public key $A: A = [b - A']$, secret key: s

→ $\left\{ \begin{array}{l} C = mt + A'^T r \\ \text{distinguishing } A \text{ from random matrix} \end{array} \right.$

$$\left\{ \begin{array}{l} d = \langle s, c \rangle \text{ mod } q \text{ mod } 2 = s(Ar + b) = (sA)r + \langle s, b \rangle = \underbrace{\langle e, r \rangle}_{-1 \leq r \leq q-1} + \underbrace{\langle s, b \rangle}_{-1 \leq b \leq q-1} \text{ mod } q \end{array} \right.$$

- Addition

$$C_1 + C_2 =$$

- Multiplication:

$$C_1 \cdot C_2$$

③ Bootstrapping

- SHE with pk_1 and sk_1
- $c_1 = \text{Enc}(m, pk_1)$
- $\widehat{sk}_1 = \text{Enc}(sk_1, pk_2)$

Recrypt

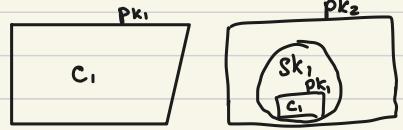
$$\bar{c}_1 = \text{Enc}(c_1, pk_2)$$

$$c = \text{Enc}(\text{Dec}(\widehat{sk}_1, \bar{c}_1), pk_2)$$

$$\text{Enc}(\text{Dec}(\text{Enc}(sk_1, pk_2), \text{Enc}(c_1, pk_2)), pk_2)$$

- Gentry's bootstrapping theorem

if a SHE scheme can evaluate its own decryption function, then its bootstrappable and can evaluate everything - with no limits



⑤ Information Centric Networking

- What if we want to ask just for data by its unique name?

① Idea

- Location → name of data

② Named Data Networking

- Every piece of data named uniquely; names are hierarchical

e.g. /ndn/ucla/videos/

- IP <=> NDN (fetch data chunks by names)
 - ↓
 - no address, only app-def namespaces
 - NDN names data instead of data location
 - Consumers fetch data instead of senders pushing packets to destinations

- NDN node's forwarding module

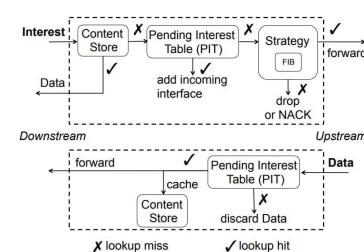


Fig. 3. Forwarding Process at an NDN Node: upstream indicates the direction of data producer and downstream is to data consumer.

- FIB - forwarding information base

- NDN for. uses the name in Interest (signed Interest)

③ DONA (Data-Oriented Network Arch.) - pull model

- organized around Principles

- flat names $\langle P \cdot L \rangle$, P - hash of principal public key, L - label unique chosen by P (hash of object)

- searching by name, Resolution Handlers (RH) name-to-data mapping, similar to DNS

- 3 packet types:

- find • data = register

④ PubSub (communication model) - push model

⑤ Named Function Networking (NFN)

- get the result, not the data

- Lambda-calculus:

- Formal system for the def. and manipulation of functions

- Effectively computable funct. definable in the lambda calculus

- Equivalence to Turing Machines

$$2. (\lambda y. [x(y z)]) (ab)$$

↑
head ↑
 body

- Call-by-value vs call-by-name vs call-by-need

$$(\lambda x \rightarrow x + x) (stu)$$

- NFN: users send block of functions(as lambda)

- network decides how to distribute comput., deliver the result, use cached

- call-by-name

- NFN exist in each node

- NFN Node Model

◦ Service Layer

- manipulates data
- fetches resources

◦ NFN Layer

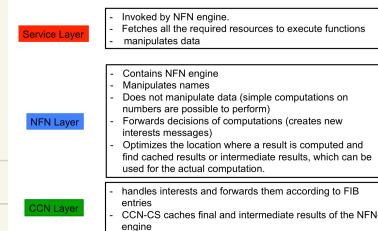
- NFN engine
- manipulate names
- forwards decisions of computation
- optimizes the locat.

(NDN - Named Data Networking)

◦ CCN Layer (Content Centric Networking)

- handles interests and forwards them according FIB entries

- CCN-CS caches final and intermediate results of the NFN-engine



⑥ Inter-Planetary File System (IPFS) (BitTorrent)

- o DS for storing and accessing files
- o content addressable
- o Peer-to-Peer
- o retrieves data from nearest nodes

⑦ Secure Scuttlebutt (SSB)

- o decentralized messaging and social networking protocol)
- o content addressable data: each piece of content is signed and append-only list of mem
- o each user hosts their own content and the content of the peers they follow

⑥ CRDT

- How to make all replica strong eventually consistent?

① Introduction

- We want to make concurrent transactions, but some of them can be conflicted.
- Database approach
- The blockchain approach
- Replication protocols
 - where
 - when
- primary copy
- update everywhere
- eager replication
- lazy replication
- with lazy update everywhere we have no consistency

2. Replicas protocol

1. Random Notepad

- Update: 1 Step: dice
 -
 - ↓ nothing
 - reconcile with replica 1-4 if not busy
 - ↓
 -
 - initiator number
 - responder number
- Write everywhere:
 - choose [1, 9]
 - pick 1 replica
 - replica replace

Number with client number

2 Novelty Notepad Protocol

- Lazy update
- Update: 1 Step: dice
 -
 - ↓ nothing
 - reconcile with replica 1-4 if not busy
 - ↓
 - deterministic
 - if tmp > both replica use most recent value
 - keep current
- Write everywhere:
 - choose [1, 9]
 - pick 1 replica
 - Client reads the current tmp of replica, tmp+1 and write new value

3. The Post-It Protocol

- Lazy update
- Update: 1 Step: dice
 -
 - ↓ nothing
- pull:
 - m > 1 you don't own
 - give [1, m-1]
- push
- Update Everywhere
 - Client picks one replica
 - [1, 9] not already in that replica. write ≥ 4 copies
 - write to replica

- consistency is stable with a non-null probability

- NO, if writes are made on most recent tmp
 - yes, most recent writes are concurrent

- Eventually consistent

- long time to stability
- which state consistent can't be predicted

- may stay inconsistent with concurrent operations

- strong eventual consistency

③ Conflict-Free Replicated Data Types (CRDTs)

- Strong eventual consistency:

1. Eventual update: if an update is applied by a correct replica, then all correct replicas
2. Convergence: any 2 correct replicas that have applied the same update are in the same state.

- Ways to design and reasoning about CRDTs:



Operation View:

$S_{i+3} = S_{j+2}$ iff any concurrent op_k, op_l are commutative



State view

- $S_{i+3} = S_{j+2}$ iff:

1. $S_{k+1} = op_k(S_k, \dots)$ implies that $S_{k+1} \sqsupseteq S_k$ (op_k is monotonic)
2. $S_{k+1} = S_k \sqcup S_k$ implies that S_{k+1} is LUB S_k, S_k

- Type of CRDTs:



- Op-based Grow-Only Set

- State: $S_0 \leftarrow \{\}$

- Merge: $S_{i+1} \leftarrow add(add(S_i, x), \dots)$

- Operations: $S_{i+1} \leftarrow add(S_i, x)$
 $S_{i+1} = S_i \cup \{x\}$

- eventually converges as long as all operations are reliably delivered to all replicas.

- State-based Grow-only set

- State: $S_0 \leftarrow \{\}$

- Merge: $S_{i+1} \leftarrow S_i \sqcup S_j$
 $S_{i+1} = S_i \cup S_j$

- Operations

$S_{i+1} \leftarrow add(S_i, x)$:

$S_{i+1} = S_i \cup \{x\}$

- eventually converges even with msg. dropped, reordered, duplicated

- requires: 1) eventual delivery of some msg.
2) transitive connectivity between all replicas

- Partial order

- way to define a relationship between elements

- $S_1 \sqsubseteq S_2 \Leftrightarrow S_1 \subseteq S_2$

- is a binary relation \sqsubseteq between $x_i, x_j \in S$ such that $x_i \sqsubseteq x_j$ is reflexive, transitive and antisymmetric.

- Monotonic Operations: $x_{i+1} \leftarrow op(x_i, \dots)$ is m iff for all values of x_i and arguments $(\dots), x_i \sqsubseteq x_{i+1}$

- Merge as Least-Upper Bound (LUB)

- at least as large as both S_i and S_j

- as small as possible

Properties:

- commutative

- associative

- idempotent

④ Designing with State-based CRDTs

- 1. Set of Possible Values
- 2. Partial order
- 3. Monotonicity
- 4. LUB (Merge)

→ SEC

① Single-Writer Grow-Only Counter

- CRDT design where only one writer that increments the counter.

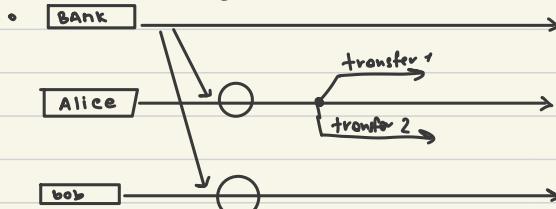
② Casual-Length Set

- adding + removing elements
- S_i - dictionary, $\{\text{element: state}\}$
 - if $S_i[k]$ is odd, k in set
 - if $S_i[k]$ is even, k is not in set
- Set of possible states: $\text{KEYS} \times [0, \infty]$
- Partial Order: $S_i \sqsubseteq S_j : \text{keys}(S_i) \subseteq \text{keys}(S_j)$ and $\forall e \in S_i[e] \leq S_j[e]$
- Monotonic operations:
 - add
 - remove
- Merge (LUB): max

⑦ Double-Spending

① Problem

- Double-Spending comes from concurrent updates from the same author



- Detection: this will be eventually detected by correct replica by merging both victims
- Prevention: restrict state changes, make them sequent.

② Append-Only Log (AOL)

- Sequence of immutable events/msm from a single processor/author.

- In terms of CRDT:

- Set of possible values: $L_i = ((L_i)_{\text{author}}, (L_i)_{\text{last}}) \in S = \mathcal{A} \times \mathcal{M}$ such that $(L_i)_{\text{last}} \leq (L_j)_{\text{last}}$

- Partial Order: $L_i \leq L_j : L_i, L_j \in S$ and $(L_i)_{\text{author}} = (L_j)_{\text{author}}$ and $(L_i)_{\text{last}}$ ancestor of $(L_j)_{\text{last}}$

- Monotonic operation:

```

 $L_{i+1} \leftarrow \text{append}(L_i, m) : L_i \in S \text{ and } m \in M$ 
 $L_{i+1} \leftarrow L_i$ 
 $\text{if } m_{\text{author}} = (L_i)_{\text{author}} \text{ and } (L_i)_{\text{last}} \text{ ancestor of } m$ 
 $\text{then}$ 
 $(L_{i+1})_{\text{last}} \leftarrow m$ 
  
```

```

 $L_{i+1} \leftarrow L_i \sqcup L_j : L_i, L_j \in S \text{ and } (L_i)_{\text{author}} = (L_j)_{\text{author}}$ 
 $(L_{i+1})_{\text{author}} \leftarrow (L_i)_{\text{author}}$ 
 $\text{if } (L_i)_{\text{last}} \text{ ancestor of } (L_j)_{\text{last}}$ 
 $\text{then}$ 
 $(L_{i+1})_{\text{last}} \leftarrow (L_j)_{\text{last}}$ 
 $\text{else if } (L_j)_{\text{last}} \text{ ancestor of } (L_i)_{\text{last}}$ 
 $(L_{i+1})_{\text{last}} \leftarrow (L_i)_{\text{last}}$ 
  
```

undefined if both lasts are concurrent

- Unique author per log in replica + no concurrency of log replicas, so we can represent the state of a replica as a vector clock
ex: (alice, 2)

③ Git Commit Graph

- in append-log there is no causal relationship between logs.
(how one entry may depend on or relate to another)

④ Secure Append-Only Log

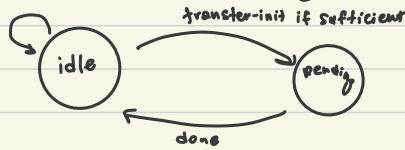
- 2P-BFT-Log

• problem: totally ordering msm from a given author, similar to an append-only log, while providing eventual consistency and fault-detection in the presence of Byzantine authors that create concurrent msm.

- 1st phase: growing
- 2nd phase: shrinking
- fork \Rightarrow double-spending, double-spending \Rightarrow was fork!

⑤ Preventing Double-Spending with Majority Ack Protocol

- Quorum
- To avoid stuck in the pending state forever



- 2PC vs Append-only logs

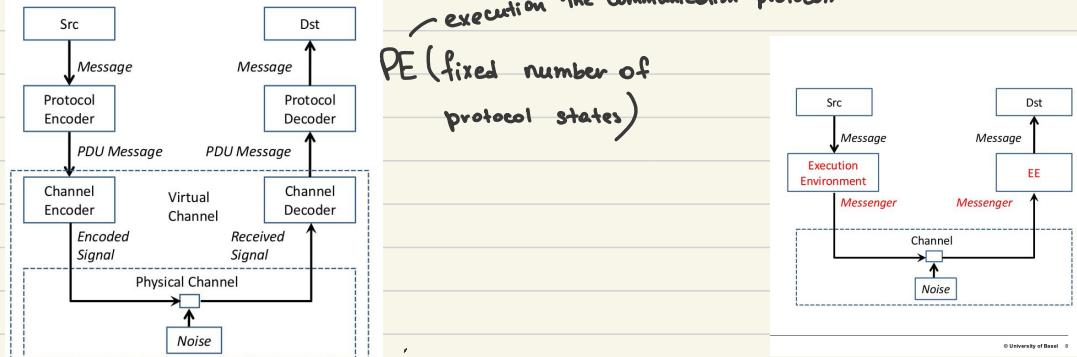
+ two phase

- diff:
1. tolerates a subset of Byzantine behaviour
 2. no abort, just ignore
 3. no ack, just tracking others
 4. non-blocking

⑧ Software-Defined Networking

① Active Networking

- Shannon's communication model
- PDU - Protocol Data Unit, PE - Protocol Entity



Idea: want generic environment to compute whatever program it receives.

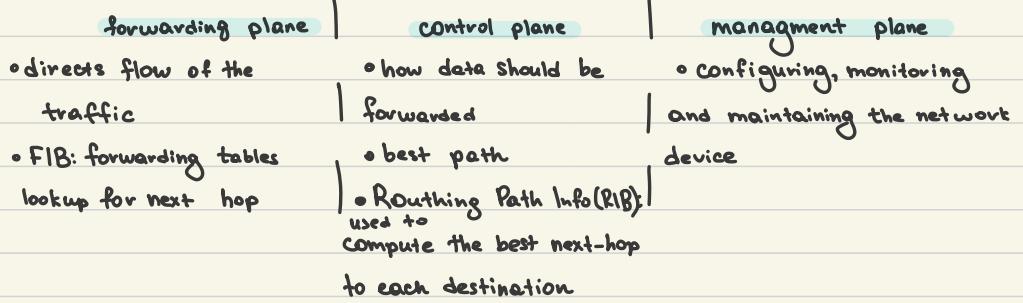
So: PE → Execution Environment

PDU → Messenger of Imperative Expression + Data

• Memes concept: Receiver

② Software Defined Networking (SDN)

- Every network device has 3 "planes" of operations



- Classic routing protocols in the Internet today

④ distributed and

scalable

reliability and fault tolerance

⑤ Most distributed

algo are more complex

to implement than their centralized counterparts

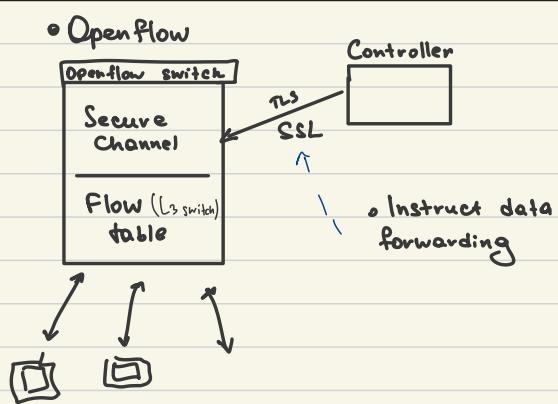
- SO: make them a little bit centralized for simplicity for programmability:

• forwarding → distributed

• controller → centralized

- We remove the control plane routing protocols from a router:

result => forwarding by arbitrary flow rule => OpenFlow switch



• Flow entry (table):

| match fields | counters | instructions |
|------------------|----------|------------------------|
| | ↓ | ↓ · APPLY |
| · IP | | ↓ · CLEAR |
| · TCP/UDP port | | ↓ · ADD |
| | | ↓ · Include (metadata) |

• NVN

(10) Bio-inspired

① Introduction

Examples: 1. Bionics 2. Embryonics 3. ANN u. ALife 5. Artificial Chemistry
6. Chemical Computing Models 7. Complex Systems 8.

② Example: The Note-and-Spread Game

③

- find max \Rightarrow chemistry

- 5 properties of reacting systems:

1. NO distinctions between passive and active

2. Objects can transform into other object without interactions

3. Objects can interact with other objects to create one or more objects

4. Objects can interact in parallel

5. can regenerate

- **[def]** AC: (S, R, A) : S-set of molecules, R-set of reactions, A-scheduling algo

- Levels of computations:

1. Symbolic Chemical Comp. (the result find in a single molecule)

2. Organizational CC (The result is find in the presence/absence of species)

3. Dynamic CC (The result in the number of instances of a species)

- Idea of Dynamic CC:

1. Info: concentration of molecules

2. State: vector of concentr. of all species

3. Computation: change in concentration

1. Law of Mass Action (LoMA):

The average reaction rate (speed) is proportional to the concentration of its reactants"

- $r = \alpha_1 S_1 + \alpha_2 S_2 + \dots + \alpha_{|S|} S_{|S|} \xrightarrow{k} \beta_1 S_1 + \dots + \beta_{|S|} S_{|S|}$

- $v = k \cdot \prod (C_{S,i})^{\alpha_i}$

C-concentration, α_i -consumed by r, β_i -created by r

- Links concentration and rates of a system

• ODEs

$$\dot{c} = \sum v, [\sum]_{S,r} = \beta_{S,r} - \alpha_{S,r}$$

2. Equilibrium in reaction systems:

1. Law of Mass Action: speed proportional concentration

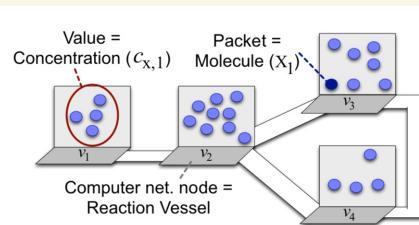
2. Flow conservation principle: concentr=const, speed the same

3. Mass conservation principle: the number of molecules is conserved in reaction loops

3. Stochastic simulation algo

1. Naive

2. Next method



④ Chemical Networking Protocols and Algorithms:

1. Distributed Artificial Chemistry

- Each species is a node in computer network

• Reaction model of chemistry: reactions consume molecules in one node and produce molecules in other nodes

• Reaction types:

- Local reactions: local consume, local produce

- Communicating/distributed reaction: local consume, local/remote produce

2. Gossip-style: Chemical Algorithms

Task: compute an average value in distributed manner

$$I. \text{Gossip Approach: } S_n = \frac{\text{Sold} + \text{Space}}{2}$$

⇒ hard protect malicious attacks, not robust

II. Chemical alternative:

Idea: Concentration represent the approx. value

- Send molecules from one node to a neighbor node

• Disperser - Chemical Protocol

• Consensus in Wireless Sensor Networks (WSNs)

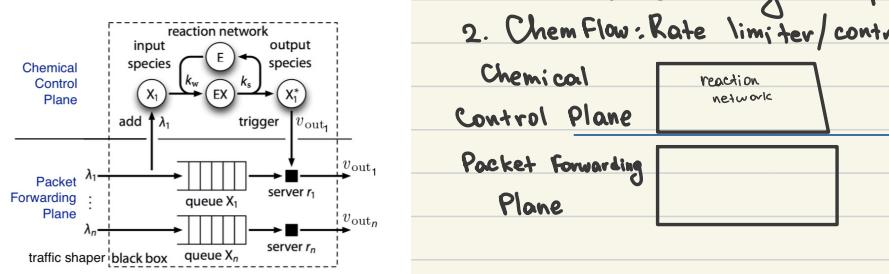
- Idea + embodiment concept

- Nodes end up in a correct state, independently of their initialization + cope with not foreseen faults.

3. Chemical Rate Control Algorithms

1. LoMA as scheduling discipline

2. ChemFlow: Rate limiter/controller



Chemical Control Plane

Packet Forwarding Plane

Reaction network

• Enzymatic model

• E and EX; $C_E + C_{EX} = \text{const} = C_0$, $V_{max} = k_{max} C_0$

• Low-pass filtering ($k_s \sim$ cut-off freq.)

• LoMA vs Enzymatic: Run-time programmability

4. Distributed Rate Controller

• Problem

- Idea: look at local load and load at peers