

① Basic abstraction

1. Process and Link

2. Properties: Civility and safety

3. Process failures (crash-fault, omissions, crash-recovery, eavesdropping, arbitrary faults) Byzantine

4. Crypto abstraction: hash, MAC, Digital signature.

5. Abstraction communication: fully-connected mesh, ring, broadcast medium, mesh of links intercon.

6. Link abstraction: (All are point-to-point link).

1. Fair-loss links (FLL₁, FLL₂, FLL₃)

Crash-stop

- 2. Stubborn Links (SL₁, SL₂)
- 3. Perfect Links (PL₁, PL₂, PL₃)

Crash-Recovery

- 4. Logged Perfect Links (LPL 1 → + Indication (stable Storage))

Byzantine

- 5. Authenticated Perfect Links (AL₁ → + Indication)

Algo: Retransmit Forever
(every step, just send again and again)
→ no duplication

Algo: Eliminate Duplicates
(differ only unique, receive everything)

Algo: Log Delivered (+ safe not delivered)

reliable delivery
indication
(Lpl, Delivered/received)

Algo: Authenticate and Filter (MAC)
Authenticity (instead other creation, check if sender correct)

1. FLL1: Fair-loss: If a correct process p infinitely often sends a message m to a correct process q , then q delivers m an infinite number of times.

FLL2: Finite duplication: If a correct process p sends a message m a finite number of times to process q , then m cannot be delivered an infinite number of times by q .

FLL3: No creation: If some process q delivers a message m with sender p , then m was previously sent to q by process p .

2. SL1: Stubborn delivery: If a correct process p sends a message m once to a correct process q , then q delivers m an infinite number of times.

SL2: No creation: If some process q delivers a message m with sender p , then m was previously sent to q by process p .

3.

PL1: Reliable delivery: If a correct process p sends a message m to a correct process q , then q eventually delivers m .

PL2: No duplication: No message is delivered by a process more than once.

PL3: No creation: If some process q delivers a message m with sender p , then m was previously sent to q by process p .

4.

LPL1: Reliable delivery: If a process that never crashes sends a message m to a correct process q , then q eventually log-delivers m .

LPL2: No duplication: No message is log-delivered by a process more than once.

LPL3: No creation: If some process q log-delivers a message m with sender p , then m was previously sent to q by process p .

5.

AL1: Reliable delivery: If a correct process sends a message m to a correct process q , then q eventually delivers m .

AL2: No duplication: No message is delivered by a correct process more than once.

AL3: Authenticity: If some correct process q delivers a message m with sender p and process p is correct, then m was previously sent to q by p .

7. Time Assumptions:

1. Synchronous System

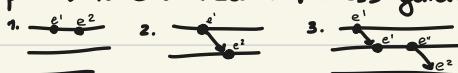
- We know: processing delay + transmission delays + local physical clock deviation
- $2 \cdot \Delta t_p + \Delta t_p \leq 2 \cdot t_{link} + \Delta t_q$

2. Asynchronous System

- logical time (clock)

- Algo to construct Logical clock

- Happen-before : (Each process generates events in a seq)



Opposite is not true :::

- Casual history

1. Each process p keeps an integer called logical clock l_p , initially 0.
2. Whenever an event occurs at process p , the logical clock l_p is incremented by one unit.
3. When a process sends a message, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.

8. Failure detector

Synchronous system → 1. Perfect Failure Detection (PFD₁, PFD₂)

→ 2. Leader Election (LE₁, LE₂)

3. Eventually-Perfect Failure Detector

- Algo: Exclude on Timeout
[Synchronous system, so timeout + heartbeat]

PFD₁: Strong completeness: Eventually, every process that crashes is permanently detected by every correct process.

PFD₂: Strong accuracy: If a process p is detected by any process, then p has crashed.

LE1: Eventual detection: Either there is no correct process, or some correct process is eventually elected as the leader.

LE2: Accuracy: If a process is leader, then all previously elected leaders have crashed.

ELD₁: Strong completeness: Eventually, every process that crashes is permanently suspected by every correct process.

ELD₂: Eventual strong accuracy: Eventually, no correct process is suspected by any correct process.

ELD₁: Eventual accuracy: There is a time after which every correct process trusts some correct process.

ELD₂: Eventual agreement: There is a time after which no two correct processes trust different correct processes.

u. Eventual Leader Detection (ELD₁, ELD₂)

g. Distributed Systems Models

	P	L	FD	
fail-stop	c-s	perfect	perfect	→ synchronous
fail-noisy	c-s	perfect	eventually perfect	→ partially synch.
fail-silent	c-s	perfect	none	→ asynchronous
fail-recovery	c-r	stubborn	eventual, perfect	
fail-silent-arbitrary	Byzantine	authentic, perfect	none	
fail-noisy-arbitrary	Byzantine	auth. perfect	Byzantine eventual leader detector.	

(2) Reliable Broadcast

1. Strongest assumptions (fail-stop: c-s, perfect, _____)

Best-Effort Broadcast

- Assumptions: all process know Π , p-t-p link between all Π , messages are unique.

Algorithm: O(N) Steps

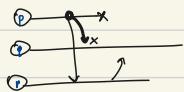
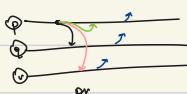
- BEB1: Validity: If a correct process broadcasts a message m , then every correct process eventually delivers m .

1. Basic Broadcast

- BEB2: No duplication: No message is delivered more than once.

- BEB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

- Sender Correct v. Some p. correct



- We need some property if sender is faulty, to make sure that all process will have the same reality. \Rightarrow inconsistencies

2. (Regular) Reliable Broadcast

Properties:

RB1: Validity: If a correct process p broadcasts a message m , then p eventually delivers m .

RB2: No duplication: No message is delivered more than once.

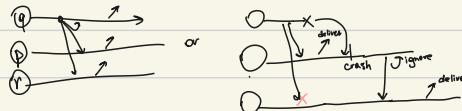
RB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

RB4: Agreement: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

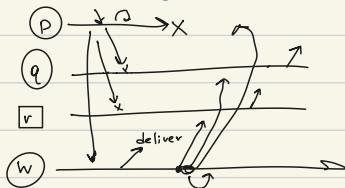
Algorithm: (- Two ways to hold „Agreement“ property: 1. relies on Perfect Failure detector 2. Take care of everything but increase complexity)
worst case

1. Fail-stop: Lazy Reliable Broadcast (Synchronous system) $O(N) + O(N^2)$ messages

- Retransmit a message only if the original sender has been detected to have crashed.



2. Fail-silent: Eager Reliable Broadcast - $O(n^2)$



3. Uniform Broadcast

(differ from reliable broadcast :
inconsistent, side-effects)



Properties:

RB1: Validity: If a correct process p broadcasts a message m , then p eventually delivers m .

RB2: No duplication: No message is delivered more than once.

RB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

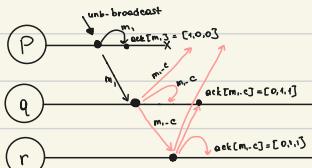
URB4: Uniform agreement: If a message m is delivered by some process (whether correct or faulty), then m is eventually delivered by every correct process.

Algorithms:

1. All-Ack URB (Perfect Failure Detector) [2 communication step, N^2 messages]

- Every process relay the message once, after they have seen it. Each process keeps a record of processes from which they

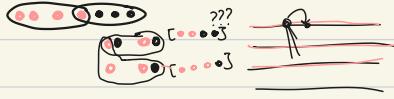
received a message. The algo uses an array $\text{ack}[m]$ gather the set of processes which seen m .



2. Majority-Ack URB (Fail-silent) [2 communication step, N^2 messages]

- Assumes a majority of correct processes i.e. $N > 2f$. So don't wait response from all process, just majority.

def Quorum - subset of processes such that any two quorums intersect on at least one correct process.



4. Probabilistic broadcast

Properties:

PB1: Probabilistic validity: There is a positive value ε such that when a correct process broadcasts a message m , the probability that every correct process eventually delivers m is at least $1 - \varepsilon$.

PB2: No duplication: No message is delivered more than once.

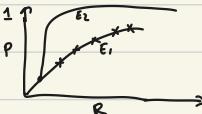
PB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

Algorithm: Eager Probabilistic Broadcast ()

- Probabilistic validity

$$E_1 = 1 - (1 - \frac{k}{N})^R$$

$$E_2 = 1 - (1 - \frac{k}{N})^{R-1} \prod_{r=0}^{R-1}$$



5. FIFO

Properties:

FRB1–FRB4: Same as properties RB1–RB4 in (regular) reliable broadcast (Module 3.2).

FRB5: *FIFO delivery*: If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

6. Casual Broadcast

Properties:

CRB1–CRB4: Same as properties RB1–RB4 in (regular) reliable broadcast (Module 3.2).

CRB5: *Causal delivery*: For any message m_1 that potentially caused a message m_2 , i.e., $m_1 \rightarrow m_2$, no process delivers m_2 unless it has already delivered m_1 .

(3) Byzantine Broadcast

① Byzantine Quorums – tolerate up to f Byzantine processes. (N process overall)
Set of more $\frac{N+g}{2}$ processes.

Q - number of nodes in a Byzantine Quorum

C correct nodes

CQ correct nodes in quorum Q

$$\#C = N-f \quad \#Q > \frac{N+f}{2} \text{ by def; } \#CQ > \frac{N+f}{2} - f = \frac{N-f}{2}$$

minimum for consistency

• Two BQ overlap in at least one correct process. consistency

- Only $N-f$ correct in the whole system.

$$\#CQ_1 + \#CQ_2 > \frac{N-f}{2} + \frac{N-f}{2} = N-f, \text{ so they intersect in 1 correct process. (Lemma)}$$

• There is at least ONE BQ with all correct. (What is f in relation to N)

$$\#C = N-f. 3Q_1 \text{ with correct p. } \#CQ_1 > \frac{N+f}{2} - f,$$

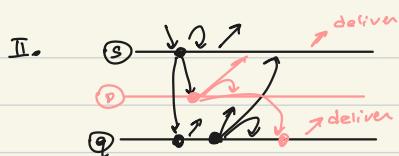
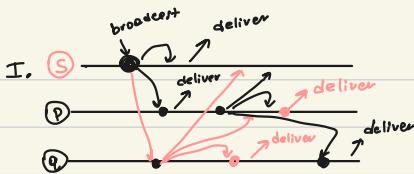
$$\#C \geq \#CQ_1 > \frac{N+f}{2}$$

$$N-f > \frac{N+f}{2}, 2N-2f > N+f, \boxed{N > 3f}$$

• Majority in a BQ has to be correct p.

② Byzantine consistent Broadcast (Fair-arbitrary model)

- Remark: No exact violation of duplication and agreement, but we have smth looks like a duplication.



- This was dubitation
 - Agreement: different delivery orderings.

Properties:

BCB1: Validity: If a correct process p broadcasts a message m , then every correct process eventually delivers m .

BCB2: No duplication: Every correct process delivers at most one message.

BCB3: Integrity: If some correct process delivers a message m with sender p and process p is correct, then m was previously broadcast by p .

BCB4: Consistency: If some correct process delivers a message m and another correct process delivers a message m' , then $m = m'$.

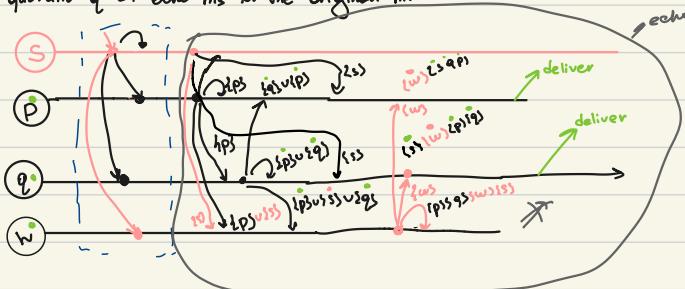
Algorithm:

hm; some processes might NOT deliver anything

1. Authenticated Echo Broadcast (Authenticated link + $N > 3f$ for consistency) [OLN²] messag[
2 step

- Two rounds of messages exchanges. No correct process delivers until they have received

a Byzantine quorum q of echo ms for the original m.



Using digital signs
we can reduce messages

③ Byzantine Reliable Broadcast (one message)

Properties:

BCB1: Validity: If a correct process p broadcasts a message m , then every correct process eventually delivers m .

BCB2: No duplication: Every correct process delivers at most one message.

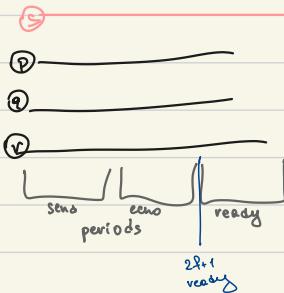
BCB3: Integrity: If some correct process delivers a message m with sender p and process p is correct, then m was previously broadcast by p .

BCB4: Consistency: If some correct process delivers a message m and another correct process delivers a message m' , then $m = m'$.

BRB5: Totality: If some message is delivered by any correct process, every correct process eventually delivers a message.

- Algorithm: Authenticated Double-Echo Broadcast (3 round: 2 echo.) $[O(N^2) + 3 \text{ com. steps}]$

- $\frac{N+p_{\text{echo}}}{2} \rightarrow \text{ready}$ [ECG]
- $\frac{p_{\text{ready}}}{2} \rightarrow \text{for you} \rightarrow \text{send ready}$ ✓
so byzantine can't cause amplification
- $2 \frac{p_{\text{ready}}}{2} \rightarrow \text{deliver}$



④ Byzantine Broadcast Channels (many messages)

1. Byzantine consist channel

Properties:

BCCH1: Validity: If a correct process p broadcasts a message m , then every correct process eventually delivers m .

BCCH2: No duplication: For every process p and label ℓ , every correct process delivers at most one message with label ℓ and sender p .

BCCH3: Integrity: If some correct process delivers a message m with sender p and process p is correct, then m was previously broadcast by p .

BCCH4: Consistency: If some correct process delivers a message m with label ℓ and sender s , and another correct process delivers a message m' with label ℓ and sender s , then $m = m'$.

2. Byzantine reliable channel

Properties:

BRCH1–BRCH3: Same as properties BCCH1–BCCH3 of Byzantine consistent channel (Module 3.13).

BRCH4: Agreement: If some correct process delivers a message m with label ℓ and sender s , then every correct process eventually delivers message m with label ℓ and sender s . *(Combine Consistency and totality)*

Algorithm: Byzantine Consistent Channel

- One broadcasting per instance (seq number for every sender + seq. of instances)

4

Consensus (crash-stop only)

- Agree on common value they initially propose.
- To require consensus, a problem must lead to negative conseq. if any of  are violated.

1. Regular Consensus

Properties:

C1: *Termination*: Every correct process eventually decides some value.

C2: *Validity*: If a process decides v , then v was proposed by some process.

C3: *Integrity*: No process decides twice.

C4: *Agreement*: No two correct processes decide differently.

Algorithm:

- Oblivious (Ignore proposals)
- Eager Best-Effort (Best-effort proposals, then decide on first proposal received)
- Waiting Best-Effort (Wait and then decide, but without Failure Detection no...)
 - Impossible to reach consensus in an asynchronous system even when only one process may crash-stop.
 - Can't have all \leq properties.
- Impossibility Result 

• Flooding Consensus (fail-stop)

- Monotonic Rounds